

Clojure

A Dynamic Programming Language for the JVM

Rich Hickey

Agenda

- Introduction
- Fundamentals
- Rationale
- Feature Tour
- Status
- Q&A

Introduction

- Who am I?
- Who are you?
 - Know/use Lisp?
 - Java?
 - ML/Haskell?

Clojure Fundamentals

- Functional
 - Mostly
- Lisp
 - *Not* CL or Scheme
- Hosted
 - On the JVM
- Supporting Concurrency
- Open Source

Rationale

- Why yet another Language/Lisp?
- No existing language has desired fundamentals
 - Functional
 - Lisp
 - JVM Hosted
 - Thread-aware

Lisp is Good

- Often emulated/pillaged, still not duplicated
- Lambda calculus yields an extremely small core
- Almost no syntax
- Core advantage still code-as-data and syntactic abstraction

What about the standard Lisps (Common Lisp and Scheme)?

- Slow/no innovation post standardization
- Core data structures mutable, not extensible
- No concurrency in specs
- Good implementations already exist for JVM (ABCCL, Kawa, SISC et al)
- Standard Lisps are their own platforms

Clojure is a Lisp not constrained by backwards compatibility

- Extends the code-as-data paradigm to maps and vectors
- Defaults to immutability
- Core data structures are extensible abstractions
- Embraces a platform (JVM)

Functional Programming is Good

- Immutable data + first-class functions
- Could always be done in Lisp, by discipline/convention
 - But if a data structure can be mutated, dangerous to presume it won't be
 - In traditional Lisp, only the list data structure is structurally recursive
- Pure functional languages tend to strongly static types
 - Not for everyone, or every task

Clojure is a functional language with a dynamic emphasis

- All data structures immutable & persistent, supporting recursion
- Dynamic typing
 - Heterogeneous collections, argument and return types
- Runtime polymorphism

Languages and Platforms

- VMs, not OSes, are the target platforms of future languages, providing:
 - Type system
 - Dynamic enforcement and safety
 - Libraries
 - Huge set of facilities
 - Memory and other resource management
 - GC is platform, not language, facility
 - Bytecode + JIT compilation

Language as platform vs. Language + platform

- Old way - each language defines its own runtime
 - GC, bytecode, type system, libraries etc
- New way (JVM, .Net)
 - Common runtime independent of language
- Platforms are dictated by clients
 - Huge investments in performance, scalability, security etc.

Java/JVM *is* language + platform

- Not the original story, but other languages for JVM always existed, now embraced by Sun
- JVM has established track record and trust level
 - Now also open source
- Interop with other code always required
 - C linkage insufficient these days
 - Ability to call/consume Java is critical
- Clojure is the language, JVM the platform

Object Orientation is Overrated

- Born of simulation, now used for everything, even when inappropriate
 - Encouraged by Java/C# in all situations, due to their lack of (idiomatic) support for anything else
- Mutable stateful objects are the new spaghetti code
 - Hard to understand, test, reason about
 - Concurrency disaster
- Inheritance is *not* the only way to do polymorphism

Polymorphism is Good

- Switch statements, structural matching etc yield brittle systems
- Polymorphism yields extensible, flexible systems
- Clojure multimethods decouple polymorphism from OO and types
- Supports multiple taxonomies
- Dispatches via static, dynamic or external properties, metadata, etc

Concurrency and state in the multi-core future

- Immutability makes much of the problem go away
 - Share freely between threads
- But changing state a reality for simulations and for in-program proxies to the outside world
- Locking is too hard to get right over and over again
- Clojure's software transactional memory does the hard part

Feature Tour

- Dynamic Development
- Lisp
- Functional Programming
- Runtime Polymorphism/Multimethods
- Concurrent Programming/STM
- JVM/Java Integration

Dynamic Development

- REPL
- Basics
 - Clojure has arbitrary precision integers, strings, ratios, doubles, characters, symbols, keywords.
- Dynamic Compilation
 - To JVM bytecode

Lisp

- Reader
 - Code-as-data
 - Extended for maps and vectors
- Macros

```
(defmacro and  
  ([ ] :t)  
  ([x] x)  
  ([x & rest] `(if ~x (and ~@rest))))
```

Functional Programming

- First-class functions

```
(defn make-adder [x]
  (let [y x]
    (fn [z] (+ y z))))
(def add2 (make-adder 2))
(add2 4)
-> 6
```

- Variable arity

```
(defn argcount
  ([] 0)
  ([x] 1)
  ([x y] 2)
  ([x y & more]
    (+ (thisfn x y) (count more))))

(argcount 1 2 3 4 5)
-> 5
```

Immutable Data Structures

- Lists, vectors, maps
- Reader, backquote support

```
(let [vec [1 2 3 4]
      map {:fred "ethel"}
      lst (list 4 3 2 1)]
  (list
    (conj vec 5)
    (assoc map :ricky "lucy")
    (conj lst 5)
    ;the originals are intact
    vec
    map
    lst))
-> ([1 2 3 4 5] {:ricky "lucy", :fred "ethel"} (5 4 3 2 1)
    [1 2 3 4] {:fred "ethel"} (4 3 2 1))
```

Persistent Data Structures

- Not DB persistence
- Old version of the collection is still available after the 'change'
- Collection maintains its performance guarantees for most operations
 - Therefore new versions are not full copies
- Hash map and vector both based upon array mapped hash tries (Bagwell)
- Sorted map is red-black tree

Extensible Abstractions

- Clojure uses Java interfaces to define its core data structures
- Supports extension to new concrete implementations
- Library functions will work with these extensions
- Big improvement vs. hardwiring a language to the concrete implementations of its data types

Example: The Seq Abstraction

- An abstraction of *cons*
 - *first* and *rest* functions
 - Immutable and persistent
 - Can be implemented lazily
- Many implementations
 - All collections provide seqs
 - Can get seqs over Java Iterables and arrays
- Library functions (map reduce filter etc) work with all seqs

Lazy Seqs

- *rest* not produced until requested
- Define your own lazy seq-producing functions using the *lazy-cons* macro
- Seqs can be used like iterators or generators of other languages

```
;the library function take
(defn take [n coll]
  (when (and (pos? n) (seq coll))
    (lazy-cons (first coll) (take (dec n) (rest coll))))))

;cycle produces an 'infinite' seq!
(take 15 (cycle [1 2 3 4]))
-> (1 2 3 4 1 2 3 4 1 2 3 4 1 2 3)
```

Metadata

- Orthogonal to the logical value of the data
- Symbols and collections support a metadata map
- Does not impact equality semantics
- Not seen in operations on the value
- Support for literal metadata in reader

```
(def vec [1 2 3])  
(def attributed-vec (with-meta vec {:source :trusted}))  
(:source ^attributed-vec)  
-> :trusted  
(eql? vec attributed-vec)  
-> :t
```

Recursive Loops

- No mutable locals in Clojure
- No tail recursion optimization in the JVM
- *recur* op does constant-space recursive looping
- Rebinds and jumps to nearest *loop* or function frame

```
(defn zipmap [keys vals]
  (loop [map {}, ks (seq keys), vs (seq vals)]
    (if (and ks vs)
      (recur (assoc map (first ks) (first vs))
              (rest ks)
              (rest vs))
      map)))
```

```
(zipmap [:a :b :c] [1 2 3])
-> {:b 2, :c 3, :a 1}
```

Polymorphism via Multimethods

- Full generalization of indirect dispatch
 - Not tied to OO or types
- Fixed dispatch function which is an arbitrary function of the arguments
- Open set of methods associated with different values of the dispatch function
- Call sequence:
 - Call dispatch function on args to get dispatch value
 - Find method associated with dispatch value
 - else call default method if present
 - else error

Example: Multimethods

```
(defmulti encounter (fn [x y] [(:Species x) (:Species y)]))
```

```
(defmethod encounter [:Bunny :Lion] [b l] :run-away)
```

```
(defmethod encounter [:Lion :Bunny] [l b] :eat)
```

```
(defmethod encounter [:Lion :Lion] [l1 l2] :fight)
```

```
(defmethod encounter [:Bunny :Bunny] [b1 b2] :mate)
```

```
(def b1 {:Species :Bunny :other :stuff})
```

```
(def b2 {:Species :Bunny :other :stuff})
```

```
(def l1 {:Species :Lion :other :stuff})
```

```
(def l2 {:Species :Lion :other :stuff})
```

```
(encounter b1 b2)
```

```
-> :mate
```

```
(encounter b1 l1)
```

```
-> :run-away
```

```
(encounter l1 b1)
```

```
-> :eat
```

```
(encounter l1 l2)
```

```
-> :fight
```

Concurrent Programming

- Immutable persistent data structures
- But, in most practical programs some state changes
- Refs
 - Sharing changes between threads
 - Transactional
- Vars
 - Isolating changes within threads

Refs and Transactions

- Software transactional memory system (STM)
- Refs can only be changed within a transaction
- All changes are Atomic and Isolated
 - Every change to Refs made within a transaction occurs or none do
 - No transaction sees the effects of any other transaction while it is running
- Transactions are speculative
 - Will be retried automatically if conflict
 - Must avoid side-effects!

Example: Refs

```
(import '(java.util.concurrent Executors))

(defn test-stm [nitems nthreads niters]
  (let [refs (map ref (replicate nitems 0))
        pool (. Executors (newFixedThreadPool nthreads))
        tasks (map (fn [t]
                      (fn []
                        (dotimes n niters
                          (sync nil
                            (dolist r refs
                              (set r (+ @r t))))))))
                    (range nthreads))]
    (. pool (invokeAll tasks))
    (. pool (shutdown))
    (map deref! refs)))

(test-stm 10 10 10000)
-> (550000 550000 550000 550000 550000 550000 550000 550000
    550000 550000)
```


Vars

- Like CL's special vars
 - dynamic scope
 - stack discipline
- Shared root binding established by *def*
 - root can be unbound
- Can be *set!* but only if first thread-locally bound using *binding* (not *let*)
- Functions stored in vars, so they too can be dynamically rebound
 - context/aspect-like idioms

JVM/Java Integration

- Consuming Java
 - Instantiate objects with *new*
 - Access members with dot operator
 - Inference + hints can avoid reflection
 - Import class names with *import*
- Extending Java
 - Implement interfaces dynamically with *implement* macro

Example: Swing App

```
(import '(javax.swing JFrame JLabel JTextField JButton)
        '(java.awt.event ActionListener)
        '(java.awt GridLayout))
(defn celsius []
  (let [frame (new JFrame "Celsius Converter")
        temp-text (new JTextField)
        celsius-label (new JLabel "Celsius")
        convert-button (new JButton "Convert")
        fahrenheit-label (new JLabel "Fahrenheit")]
    (. convert-button
      (addActionListener
        (implement [ActionListener]
          (actionPerformed [evt]
            (let [c (. Double (parseDouble (. temp-text (getText)))]
              (. fahrenheit-label
                (setText (strcat (+ 32 (* 1.8 c)) " Fahrenheit")))))))))
    (doto frame
      (setLayout (new GridLayout 2 2 3 3))
      (add temp-text) (add celsius-label)
      (add convert-button) (add fahrenheit-label)
      (setSize 300 80) (setVisible :t))))
(celsius)
```

JVM/Java Integration

- Clojure strings are Java Strings
- Clojure collections implement (read-only portion of) Collection interface
- Clojure fns are Callable
- Symbols and Keywords are Comparable
- Emits standard JVM debug info

Status

- Currently alpha
 - Download available
 - Everything I've talked about works
 - Some names might change
- Decent docs on web site
 - More coming
- Debugging with JSwat works pretty well
 - Breakpoints/Single-stepping
 - Locals

To Do

- More library functions
- Support for extending classes
- Editor/IDE support
 - Syntax highlighting, paren matching etc
 - Debugging
 - Java IDEs don't fully support debugging non-Java
- Fit/finish, performance tuning

More info

- Main site:
 - <http://clojure.sourceforge.net/>
- Discussion group:
 - <http://groups.google.com/group/clojure>

Thanks!

Questions?